

# cnn-dl

April 17, 2023

```
[18]: import tensorflow as tf
      from tensorflow import keras
      from tensorflow.keras import layers
      from tensorflow.keras.datasets import cifar10
      from tensorflow.keras.utils import to_categorical
      from tensorflow.keras.callbacks import EarlyStopping
```

```
[3]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 3s 0us/step

```
[4]: x_train = x_train.astype("float32") / 255.0
      x_test = x_test.astype("float32") / 255.0
```

```
[5]: y_train = to_categorical(y_train)
      y_test = to_categorical(y_test)
```

```
[6]: model = keras.Sequential(
      [
          layers.Conv2D(32, kernel_size=(3, 3), activation="relu",
↪input_shape=(32, 32, 3)),
          layers.MaxPooling2D(pool_size=(2, 2)),
          layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
          layers.MaxPooling2D(pool_size=(2, 2)),
          layers.Flatten(),
          layers.Dense(10, activation="softmax"),
      ]
      )
```

## 0.0.1 1)Method

```
[ ]: sgd = keras.optimizers.SGD(lr=0.01)
      model.compile(loss="categorical_crossentropy", optimizer=sgd,
↪metrics=["accuracy"])
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning\_rate`  
or use the legacy optimizer, e.g.,`tf.keras.optimizers.legacy.SGD`.

```
[ ]: model.fit(x_train, y_train, batch_size=64, epochs=15, validation_data=(x_test, y_test))
```

```
Epoch 1/15
782/782 [=====] - 109s 89ms/step - loss: 0.5534 -
accuracy: 0.8155 - val_loss: 1.0158 - val_accuracy: 0.7016
Epoch 2/15
782/782 [=====] - 71s 91ms/step - loss: 0.5105 -
accuracy: 0.8251 - val_loss: 1.4288 - val_accuracy: 0.6284
Epoch 3/15
782/782 [=====] - 72s 92ms/step - loss: 0.5066 -
accuracy: 0.8282 - val_loss: 1.1546 - val_accuracy: 0.6719
Epoch 4/15
782/782 [=====] - 71s 91ms/step - loss: 0.4943 -
accuracy: 0.8334 - val_loss: 1.0403 - val_accuracy: 0.7066
Epoch 5/15
782/782 [=====] - 70s 90ms/step - loss: 0.4837 -
accuracy: 0.8348 - val_loss: 1.0528 - val_accuracy: 0.7006
Epoch 6/15
782/782 [=====] - 71s 91ms/step - loss: 0.4799 -
accuracy: 0.8369 - val_loss: 0.9942 - val_accuracy: 0.7152
Epoch 7/15
782/782 [=====] - 71s 91ms/step - loss: 0.4718 -
accuracy: 0.8390 - val_loss: 1.0083 - val_accuracy: 0.7145
Epoch 8/15
782/782 [=====] - 74s 94ms/step - loss: 0.4720 -
accuracy: 0.8400 - val_loss: 1.2644 - val_accuracy: 0.6668
Epoch 9/15
782/782 [=====] - 70s 89ms/step - loss: 0.4637 -
accuracy: 0.8429 - val_loss: 1.0594 - val_accuracy: 0.6983
Epoch 10/15
782/782 [=====] - 71s 91ms/step - loss: 0.4628 -
accuracy: 0.8411 - val_loss: 1.0478 - val_accuracy: 0.7036
Epoch 11/15
782/782 [=====] - 71s 91ms/step - loss: 0.4562 -
accuracy: 0.8442 - val_loss: 1.0478 - val_accuracy: 0.7085
Epoch 12/15
782/782 [=====] - 69s 88ms/step - loss: 0.4535 -
accuracy: 0.8442 - val_loss: 1.0413 - val_accuracy: 0.7073
Epoch 13/15
782/782 [=====] - 71s 91ms/step - loss: 0.4512 -
accuracy: 0.8467 - val_loss: 1.0796 - val_accuracy: 0.6994
Epoch 14/15
782/782 [=====] - 69s 89ms/step - loss: 0.4468 -
accuracy: 0.8460 - val_loss: 1.0542 - val_accuracy: 0.7022
Epoch 15/15
782/782 [=====] - 71s 91ms/step - loss: 0.4457 -
```

accuracy: 0.8477 - val\_loss: 1.2613 - val\_accuracy: 0.6679

```
[ ]: <keras.callbacks.History at 0x7f32507626a0>
```

```
[ ]: test_loss, test_acc = model.evaluate(x_test, y_test)
     print(f"Test accuracy: {test_acc}")
```

313/313 [=====] - 4s 13ms/step - loss: 1.2613 -  
accuracy: 0.6679  
Test accuracy: 0.667900025844574

# 1 Finetune the hyperparameters to improve the performance of the CNN

## 1.0.1 2)Method

```
[7]: model.compile(loss="categorical_crossentropy", optimizer="adam",  
    ↪ metrics=["accuracy"])
```

```
[8]: model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test,  
    ↪ y_test))
```

Epoch 1/10  
782/782 [=====] - 71s 90ms/step - loss: 1.5390 -  
accuracy: 0.4511 - val\_loss: 1.2891 - val\_accuracy: 0.5381  
Epoch 2/10  
782/782 [=====] - 73s 94ms/step - loss: 1.2004 -  
accuracy: 0.5808 - val\_loss: 1.1184 - val\_accuracy: 0.6150  
Epoch 3/10  
782/782 [=====] - 71s 91ms/step - loss: 1.0717 -  
accuracy: 0.6318 - val\_loss: 1.0415 - val\_accuracy: 0.6450  
Epoch 4/10  
782/782 [=====] - 71s 91ms/step - loss: 0.9940 -  
accuracy: 0.6590 - val\_loss: 1.0068 - val\_accuracy: 0.6544  
Epoch 5/10  
782/782 [=====] - 70s 89ms/step - loss: 0.9354 -  
accuracy: 0.6802 - val\_loss: 0.9966 - val\_accuracy: 0.6601  
Epoch 6/10  
782/782 [=====] - 69s 88ms/step - loss: 0.8908 -  
accuracy: 0.6948 - val\_loss: 0.9450 - val\_accuracy: 0.6745  
Epoch 7/10  
782/782 [=====] - 72s 92ms/step - loss: 0.8478 -  
accuracy: 0.7116 - val\_loss: 0.9646 - val\_accuracy: 0.6751  
Epoch 8/10  
782/782 [=====] - 71s 91ms/step - loss: 0.8150 -  
accuracy: 0.7216 - val\_loss: 0.9199 - val\_accuracy: 0.6901  
Epoch 9/10

```
782/782 [=====] - 69s 89ms/step - loss: 0.7828 -  
accuracy: 0.7318 - val_loss: 0.9020 - val_accuracy: 0.6968  
Epoch 10/10  
782/782 [=====] - 71s 91ms/step - loss: 0.7576 -  
accuracy: 0.7376 - val_loss: 0.8907 - val_accuracy: 0.6960
```

```
[8]: <keras.callbacks.History at 0x7f32508f4700>
```

```
[9]: test_loss, test_acc = model.evaluate(x_test, y_test)  
print(f"Test accuracy: {test_acc}")
```

```
313/313 [=====] - 4s 13ms/step - loss: 0.8907 -  
accuracy: 0.6960  
Test accuracy: 0.6959999799728394
```

### 1.0.2 3)Method

```
[ ]: early_stop = EarlyStopping(monitor='val_loss', patience=3)
```

```
[ ]: model.fit(x_train, y_train, batch_size=64, epochs=100, validation_data=(x_test, y_test),  
            callbacks=[early_stop])
```

```
Epoch 1/100  
782/782 [=====] - 72s 92ms/step - loss: 0.1145 -  
accuracy: 0.8063 - val_loss: 0.1603 - val_accuracy: 0.7100  
Epoch 2/100  
782/782 [=====] - 71s 91ms/step - loss: 0.1134 -  
accuracy: 0.8078 - val_loss: 0.1619 - val_accuracy: 0.7072  
Epoch 3/100  
782/782 [=====] - 69s 89ms/step - loss: 0.1124 -  
accuracy: 0.8100 - val_loss: 0.1592 - val_accuracy: 0.7181  
Epoch 4/100  
782/782 [=====] - 73s 93ms/step - loss: 0.1106 -  
accuracy: 0.8140 - val_loss: 0.1586 - val_accuracy: 0.7175  
Epoch 5/100  
782/782 [=====] - 72s 92ms/step - loss: 0.1097 -  
accuracy: 0.8167 - val_loss: 0.1599 - val_accuracy: 0.7152  
Epoch 6/100  
782/782 [=====] - 72s 92ms/step - loss: 0.1087 -  
accuracy: 0.8194 - val_loss: 0.1624 - val_accuracy: 0.7123  
Epoch 7/100  
782/782 [=====] - 71s 91ms/step - loss: 0.1076 -  
accuracy: 0.8194 - val_loss: 0.1618 - val_accuracy: 0.7153
```

```
[ ]: <keras.callbacks.History at 0x7f32506070a0>
```

```
[ ]: test_loss, test_acc = model.evaluate(x_test, y_test)  
print(f"Test accuracy: {test_acc}")
```

```
313/313 [=====] - 5s 15ms/step - loss: 0.1618 -  
accuracy: 0.7153  
Test accuracy: 0.7153000235557556
```

### 1.0.3 4)Method

```
[ ]: model.compile(loss="binary_crossentropy", optimizer="adam",  
    ↪metrics=["accuracy"])
```

```
[ ]: model.fit(x_train, y_train, batch_size=128, epochs=15, validation_data=(x_test,  
    ↪y_test))
```

Epoch 1/15

```
391/391 [=====] - 68s 174ms/step - loss: 0.1393 -  
accuracy: 0.7592 - val_loss: 0.1613 - val_accuracy: 0.7025
```

Epoch 2/15

```
391/391 [=====] - 69s 178ms/step - loss: 0.1359 -  
accuracy: 0.7648 - val_loss: 0.1574 - val_accuracy: 0.7089
```

Epoch 3/15

```
391/391 [=====] - 67s 171ms/step - loss: 0.1336 -  
accuracy: 0.7678 - val_loss: 0.1571 - val_accuracy: 0.7064
```

Epoch 4/15

```
391/391 [=====] - 68s 175ms/step - loss: 0.1314 -  
accuracy: 0.7716 - val_loss: 0.1557 - val_accuracy: 0.7092
```

Epoch 5/15

```
391/391 [=====] - 67s 172ms/step - loss: 0.1291 -  
accuracy: 0.7767 - val_loss: 0.1567 - val_accuracy: 0.7080
```

Epoch 6/15

```
391/391 [=====] - 69s 177ms/step - loss: 0.1280 -  
accuracy: 0.7798 - val_loss: 0.1539 - val_accuracy: 0.7128
```

Epoch 7/15

```
391/391 [=====] - 67s 171ms/step - loss: 0.1259 -  
accuracy: 0.7837 - val_loss: 0.1543 - val_accuracy: 0.7148
```

Epoch 8/15

```
391/391 [=====] - 74s 190ms/step - loss: 0.1242 -  
accuracy: 0.7867 - val_loss: 0.1547 - val_accuracy: 0.7168
```

Epoch 9/15

```
391/391 [=====] - 68s 175ms/step - loss: 0.1227 -  
accuracy: 0.7891 - val_loss: 0.1531 - val_accuracy: 0.7183
```

Epoch 10/15

```
391/391 [=====] - 68s 174ms/step - loss: 0.1210 -  
accuracy: 0.7919 - val_loss: 0.1553 - val_accuracy: 0.7126
```

Epoch 11/15

```
391/391 [=====] - 69s 175ms/step - loss: 0.1196 -  
accuracy: 0.7949 - val_loss: 0.1536 - val_accuracy: 0.7195
```

Epoch 12/15

```
391/391 [=====] - 71s 182ms/step - loss: 0.1184 -  
accuracy: 0.7968 - val_loss: 0.1543 - val_accuracy: 0.7148
```

```
Epoch 13/15
391/391 [=====] - 71s 182ms/step - loss: 0.1174 -
accuracy: 0.8014 - val_loss: 0.1542 - val_accuracy: 0.7183
Epoch 14/15
391/391 [=====] - 67s 170ms/step - loss: 0.1159 -
accuracy: 0.8040 - val_loss: 0.1558 - val_accuracy: 0.7158
Epoch 15/15
391/391 [=====] - 67s 171ms/step - loss: 0.1150 -
accuracy: 0.8058 - val_loss: 0.1553 - val_accuracy: 0.7194
```

```
[ ]: <keras.callbacks.History at 0x7f32506daee0>
```

```
[ ]: test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test accuracy: {test_acc}")
```

```
313/313 [=====] - 5s 17ms/step - loss: 0.1553 -
accuracy: 0.7194
Test accuracy: 0.7193999886512756
```

After applying multiple method the last method has highest accuracy so we consider this model compiler and fit for training